

UNIVERSITY OF
COPENHAGEN



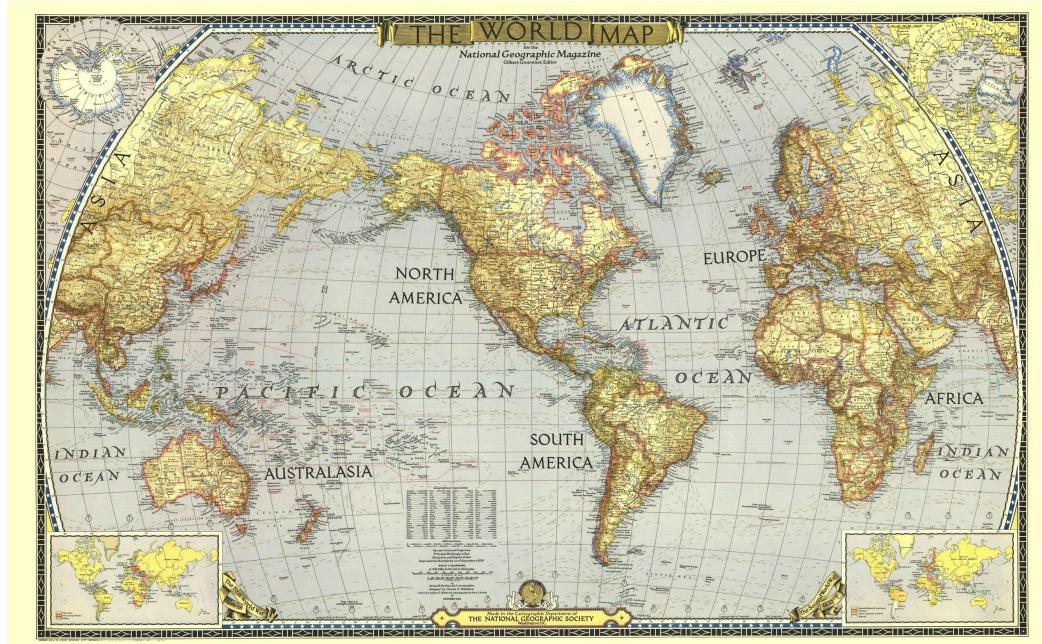
Declarative Cartography

In-Database Map Generalization of Spatial Datasets

P.K. Kefaloukos, M. Vaz Salles, M. Zachariasen

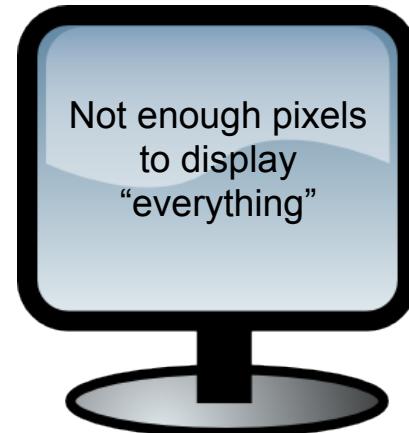
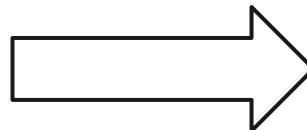
Creating zoomable maps

- Maps should be
- Legible
 - Representative
 - Useful
 - Zoomable



Basic challenge

Adapting data to scale of visualization medium



Not enough pixels
to display
“everything”

Basic challenge

We have to *choose*:



(1) *What to display?*



(2) *How to display it?*

Alternatives

Selection



Aggregation



Alternatives

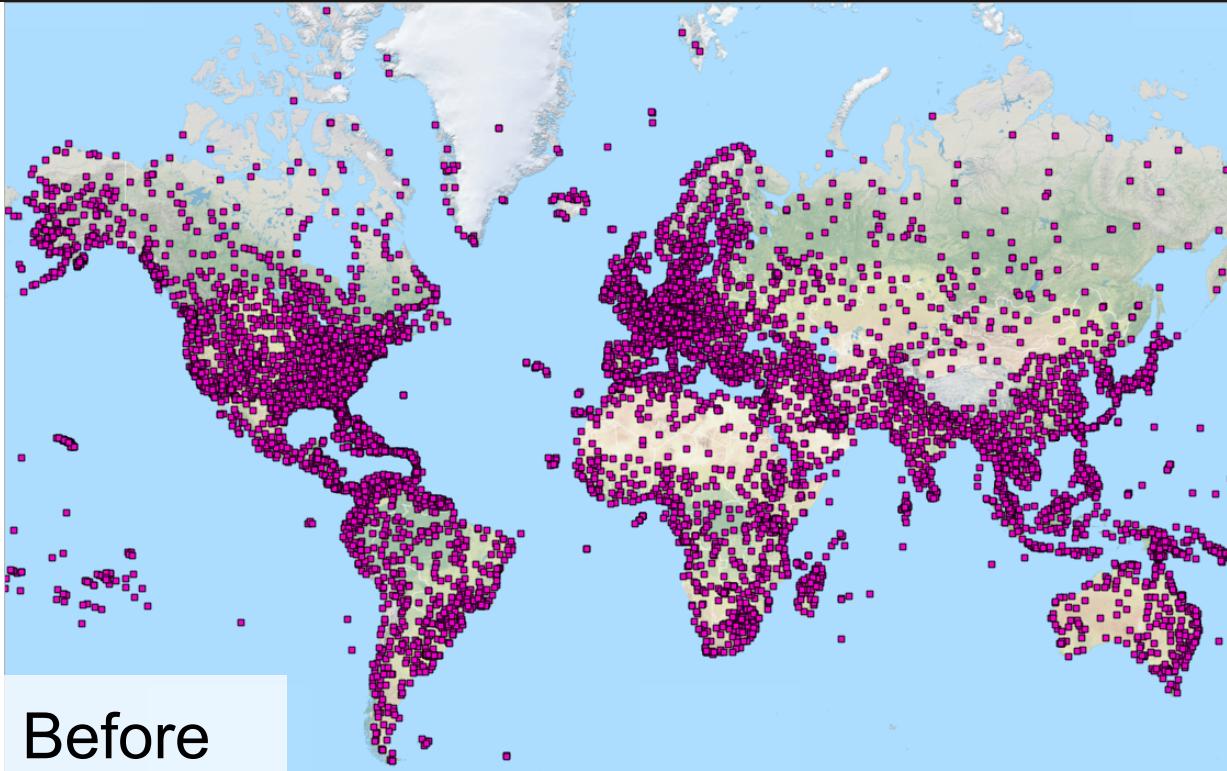
Selection (our work)



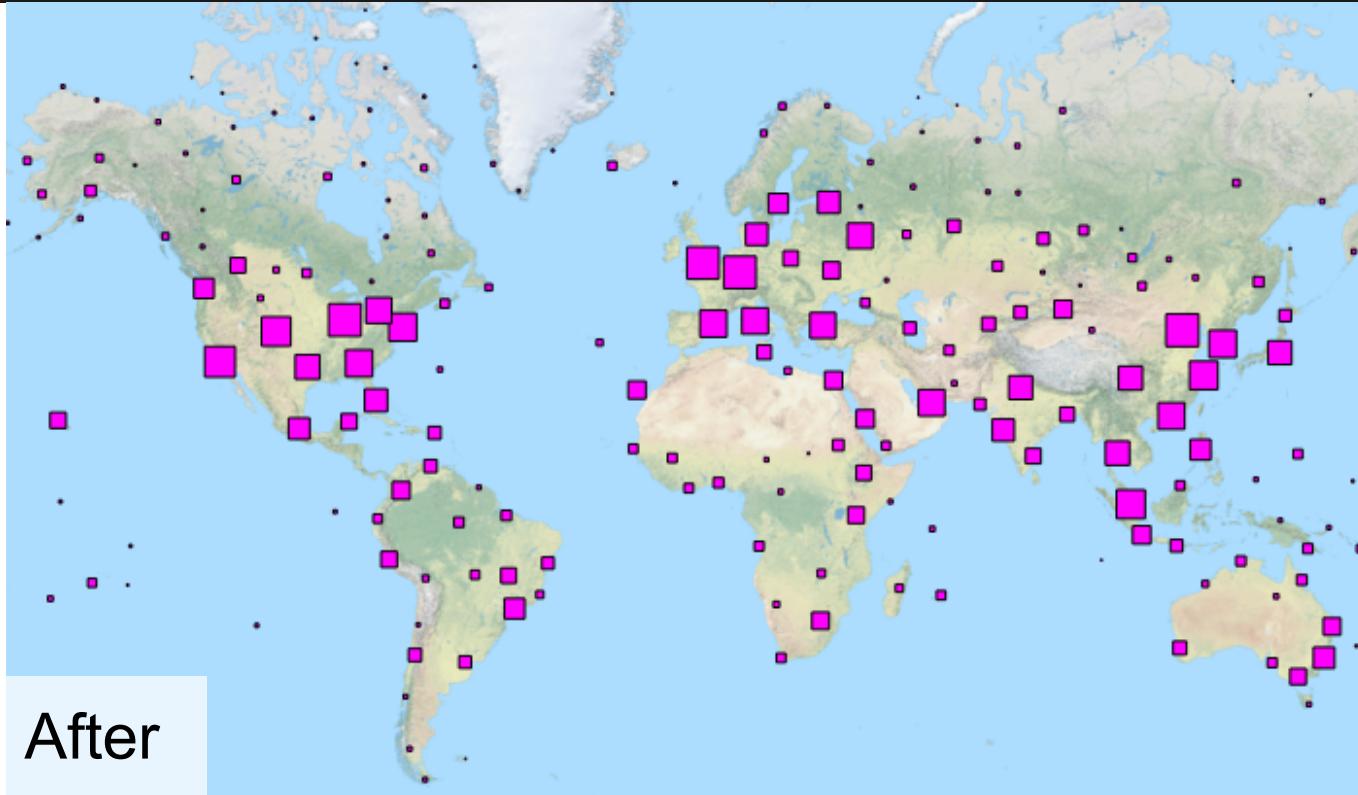
Aggregation



Selecting airports (example)



Selecting airports (example)

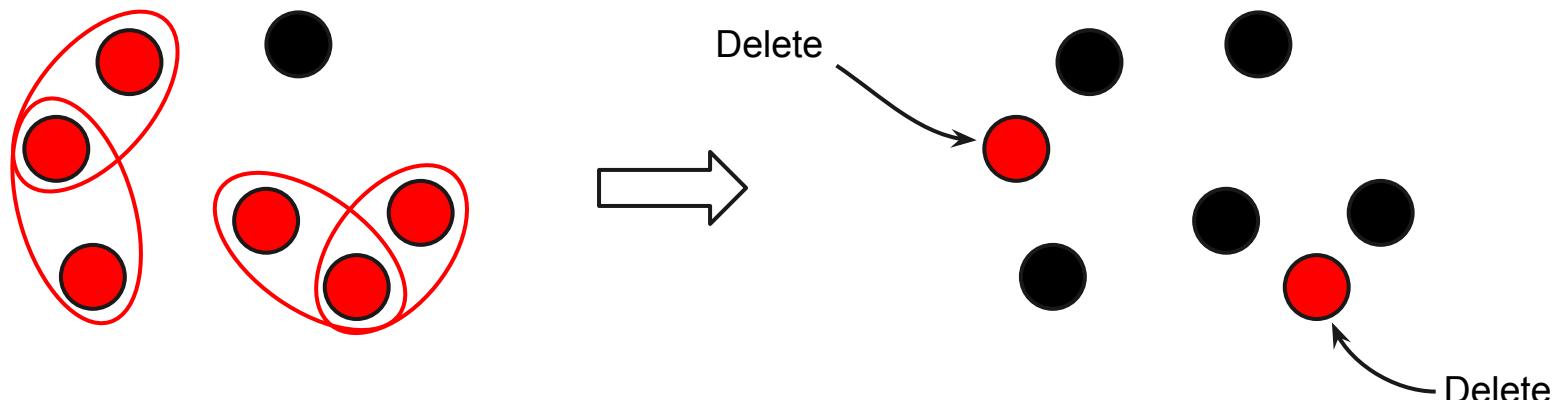


Selecting airports (example)



1: Cartographic constraints

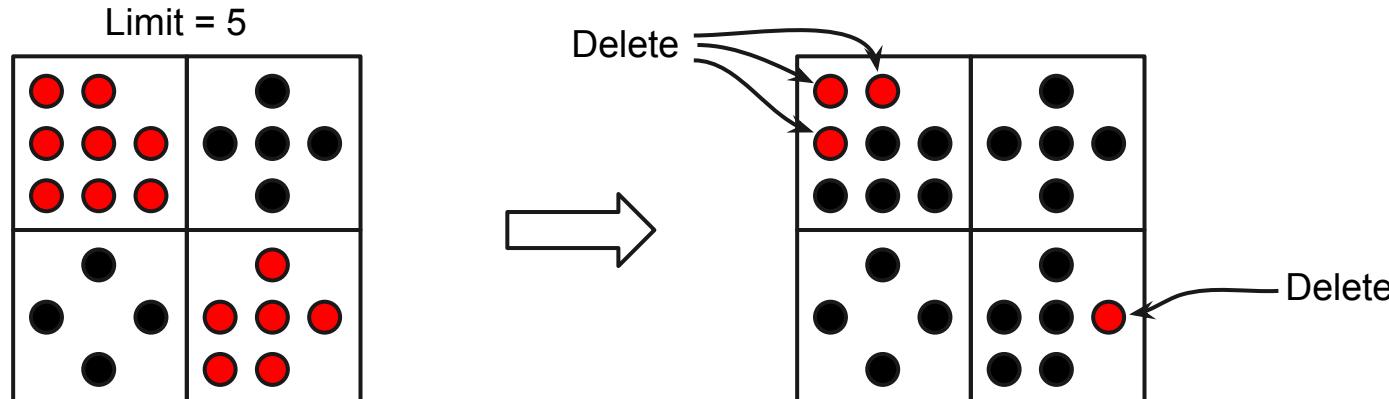
Proximity constraint: minimum distance between records
(measured in pixels on screen)



Conflicts for proximity constraint

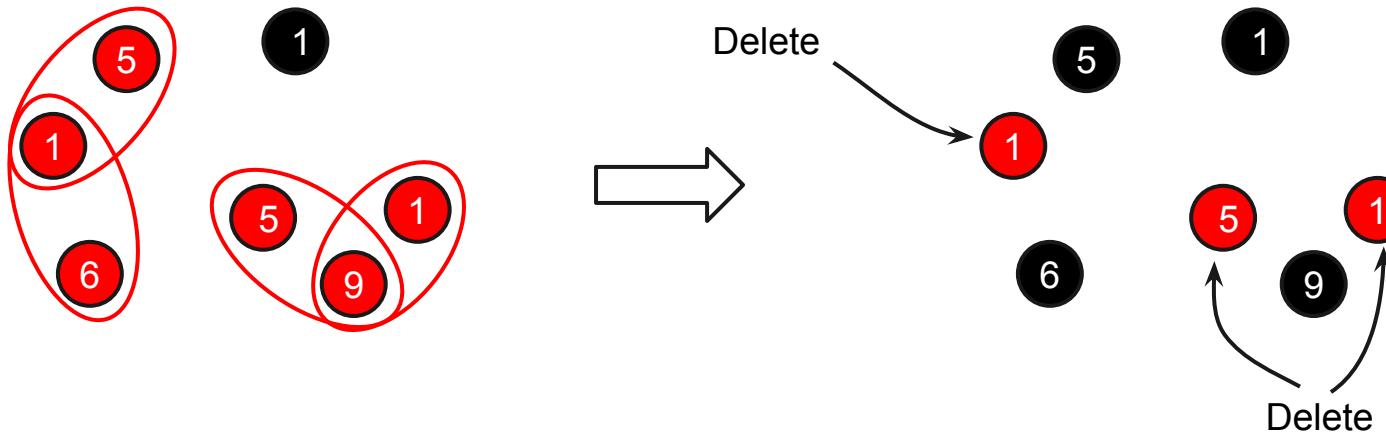
1: Cartographic constraints

Visibility: maximum records per unit area (within a map “tile”)



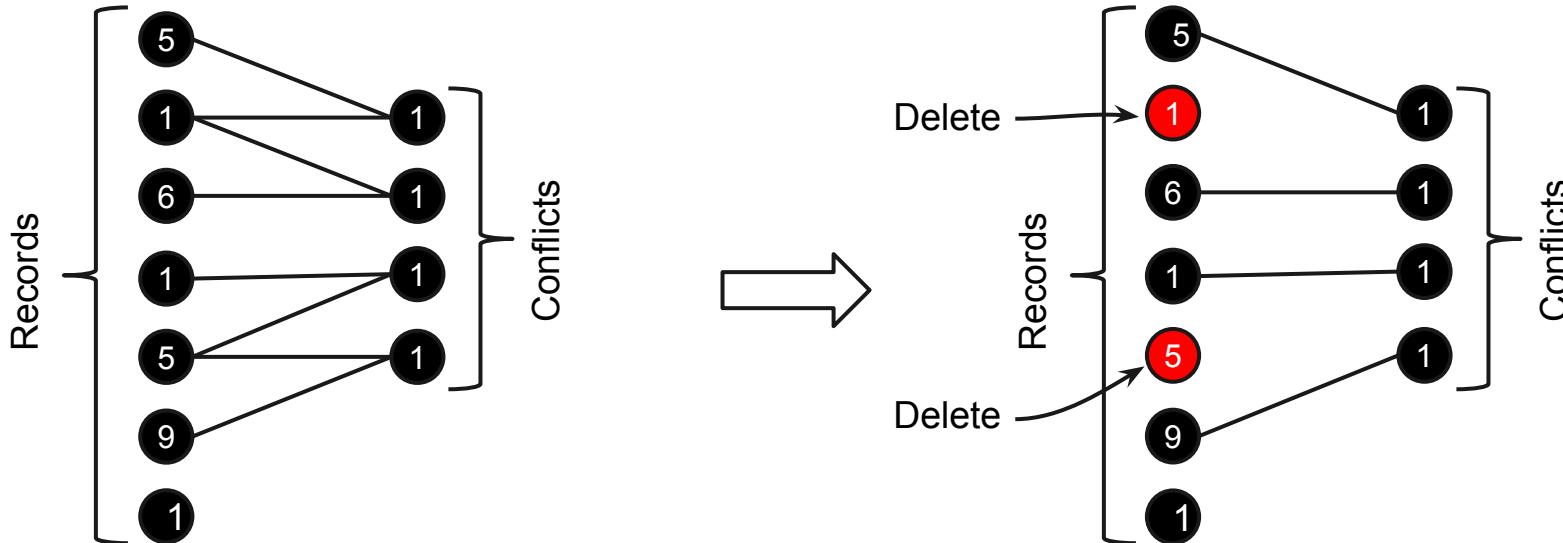
Conflicts for visibility constraint

2: Record weights



Adding weights may change selection

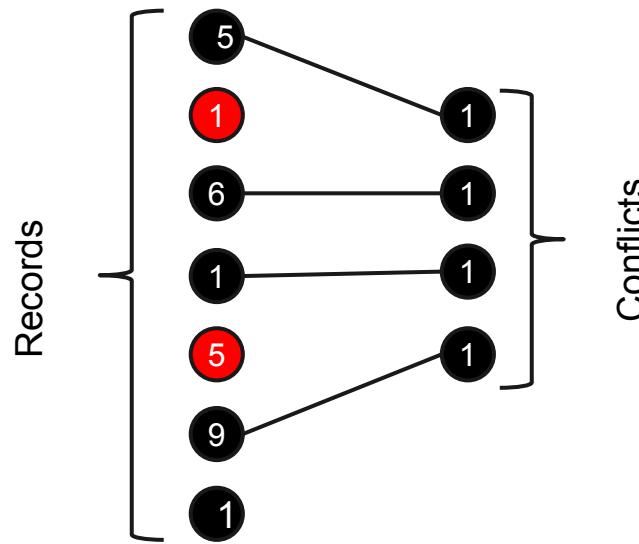
Optimization problem



Problem: delete minimum weight cover

Optimization problem

Set multicover problem
(NP-hard)



User control



Creating maps is the job of data-journalists, bloggers, high-level programmers...

... not mathematicians

Declarative Cartography



GENERALIZE

airports **TO** airports2

WITH ID airport_id

WITH GEOMETRY wkb_geometry

AT 18 ZOOM LEVELS

WEIGH BY

num_departures

SUBJECT TO

proximity 10

visibility 16

Declarative Cartography

Cartographic Visualization Language (CVL)

- Pronounced “Civil”
- State high-level goals:
 - Number of zoom-levels?
(AT 18 ZOOM LEVELS)
 - Importance of records?
(WEIGH BY)
 - Visual constraints?
(SUBJECT TO)
- Compiles to SQL

GENERALIZE

airports **TO** airports2

WITH ID airport_id

WITH GEOMETRY wkb_geometry

AT 18 ZOOM LEVELS

WEIGH BY

num_departures

SUBJECT TO

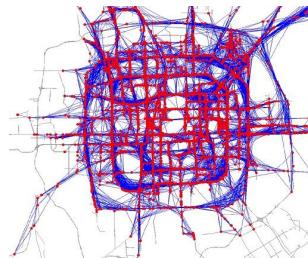
proximity 10

visibility 16

Declarative Cartography

Input dataset, e.g:

- Points of interest (points)
- Running routes (lines)
- Animal territories (polygons)



GENERALIZE

airports TO airports2

WITH ID airport_id

WITH GEOMETRY wkb_geometry

AT 18 ZOOM LEVELS

WEIGH BY

num_departures

SUBJECT TO

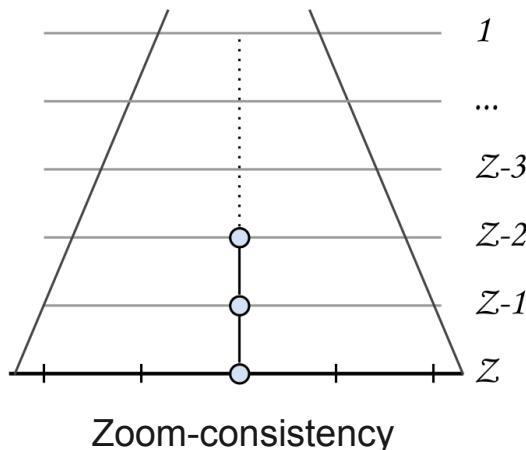
proximity 10

visibility 16

Declarative Cartography

Output dataset

- Output = input made “zoomable”



GENERALIZE

```
airports TO airports2
```

```
WITH ID airport_id
```

```
WITH GEOMETRY wkb_geometry
```

```
AT 18 ZOOM LEVELS
```

WEIGH BY

```
num_departures
```

SUBJECT TO

```
proximity 10
```

```
visibility 16
```

Declarative Cartography

Weights

- Reuse SQL expressions as sub-language
- Evaluated for each row (once)

Anything that is a number

- `random()`
- `x + y`
- `max (x, y)`

GENERALIZE

```
airports TO airports2  
WITH ID airport_id  
WITH GEOMETRY wkb_geometry  
AT 18 ZOOM LEVELS
```

WEIGH BY

```
num_departures  
SUBJECT TO  
proximity 10  
visibility 16
```

Declarative Cartography

Constraints

- Referenced by name
- Parameterized

GENERALIZE

```
airports TO airports2
```

```
WITH ID airport_id
```

```
WITH GEOMETRY wkb_geometry
```

```
AT 18 ZOOM LEVELS
```

WEIGH BY

```
num_departures
```

SUBJECT TO

```
proximity 10
```

```
visibility 16
```

Declarative Cartography

Constraints

- Referenced by name
- Parameterized

How are constraints defined?

- Let's look at proximity next

GENERALIZE

airports TO airports2

WITH ID airport_id

WITH GEOMETRY wkb_geometry

AT 18 ZOOM LEVELS

WEIGH BY

num_departures

SUBJECT TO

proximity 10

visibility 16

Declarative Cartography

General syntax

- **CREATE CONSTRAINT**: give the constraint a name
- **NOT EXISTS**: define what a conflict is (select cid, rid)
- **RESOLVE IF DELETE**: define how to resolve a conflict (number of records to delete for conflict)

```
CREATE CONSTRAINT proximity AS
  NOT EXISTS (
    /* select cid, ric ... */
  )
  RESOLVE cid IF DELETE (
    /* scalar query or integer */
  )
```

Declarative Cartography

General syntax

- **CREATE CONSTRAINT**: give the constraint a name
- **NOT EXISTS**: define what a conflict is (`select cid, rid`)
- **RESOLVE IF DELETE**: define how to resolve a conflict (number of records to delete for conflict)

```
...
NOT EXISTS (
    SELECT
        l.{rid} || r.{rid} AS cid,
        Unnest(array[l.{rid}, r.{rid}]) AS rid
    FROM
        {level_view} l JOIN {level_view} r
    ON
        ST_Distance(
            l.{geom}, r.{geom}) <
            CVL_Resolution({z}, 256) * {param_1}
)
...
...
```

Declarative Cartography

General syntax

- **CREATE CONSTRAINT**: give the constraint a name
- **NOT EXISTS**: define what a conflict is (select cid, rid)
- **RESOLVE IF DELETE**: define how to resolve a conflict (number of records to delete for conflict)

...

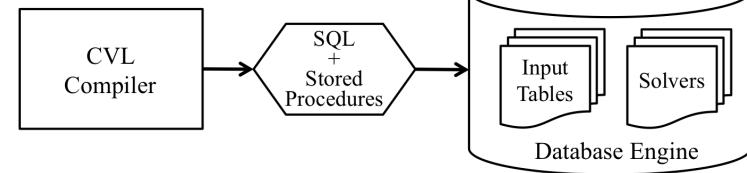
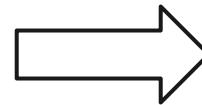
RESOLVE cid IF DELETE (

1

)

Compiling CVL into SQL

```
CREATE CONSTRAINT proximity
AS
NOT
CREATE CONSTRAINT visibility
AS
)
NOT GENERALIZE
    airports TO airports2
    WITH ID airport_id
    WITH GEOMETRY wkb_geometry
    AT 18 ZOOM LEVELS
)
RESOURCES
)
RESOURCES
)
WEIGH BY
    num_departures
SUBJECT TO
    proximity 10
    visibility 16
```



Using databases

- Spatial data typically managed in DBs
 - Option (a): Take data out of DB, process it, put it back
 - Option (b): Leverage database theory and technology to compute generalization *inside* DB (in-situ)
- We argue for option (b)
 - DBs are scalable and programmable
- However, DBs are not easy to program
(unless you use CVL!)

Results

Tested two algorithms

- Static Greedy Algorithm (SGA):
 - sorting based heuristic, implemented in SQL
- Linear Programming Greedy Algorithm (LPGA):
 - linear programming approximation algorithm,
implemented in PLPython and external LP solver

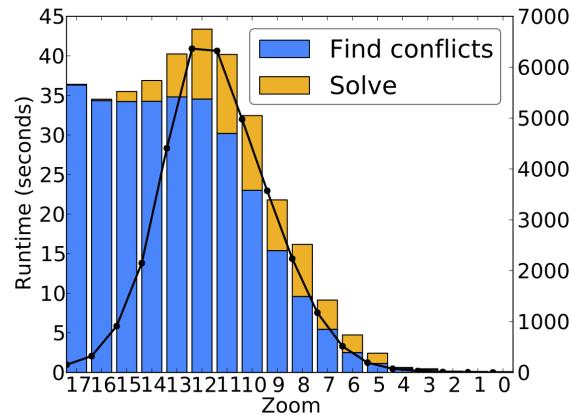
Results

Tested with one RDBMS

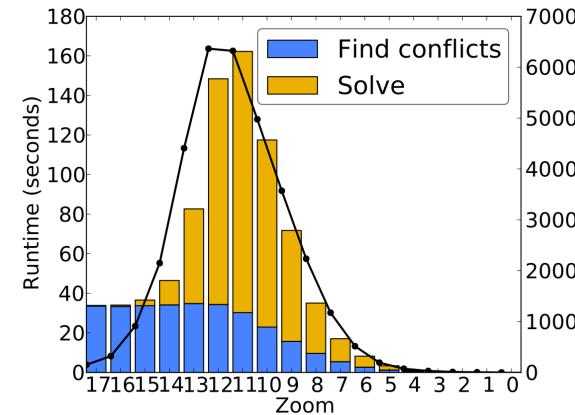
- PostgreSQL + PostGIS + Python + CVXOPT
 - Could be any database, e.g. a parallel one!

Results for points

- 500K points, tourist POI from OpenStreetMap



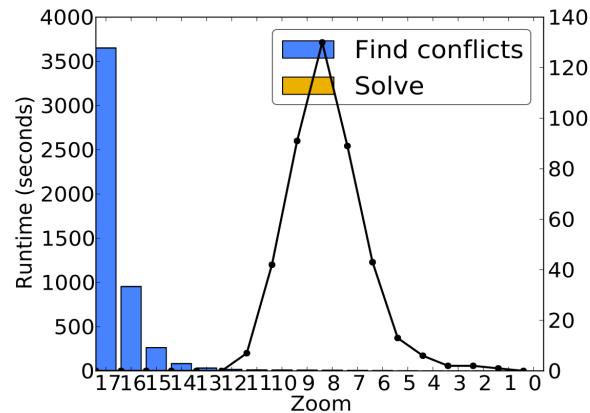
SGA + Visibility constraint



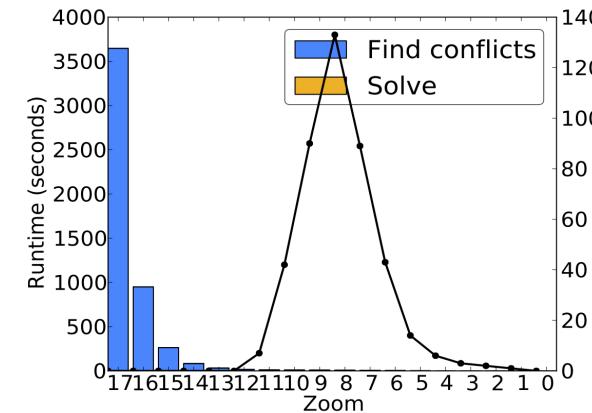
LPGA + Visibility constraint

Results for complex shapes

- 4K lines/4M points, rivers from OpenStreetMap



SGA + Visibility constraint



LPGA + Visibility constraint

Conclusion

- Core idea
 - Compute: *what* data to display on a zoomable map
 - Delay decision on *how* to display it
- Conclusion on results
 - Overall: Running time dominated by computing joins when using efficient heuristic for selection
 - Points: SQL-based algorithms perform best
 - Complex shapes: Joins for complex shapes run *much* slower than joins for points

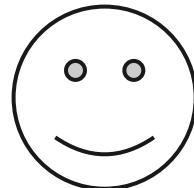
Conclusion

- Pros:
 - Domain Specific Languages make databases easier to program
 - Database engines can compute a concrete selection that matches a high-level goal expressed in DSL
 - Databases can work on data in-situ
- Cons:
 - Running time must improve
 - Selection dependent on algorithm

Conclusion

- Future work

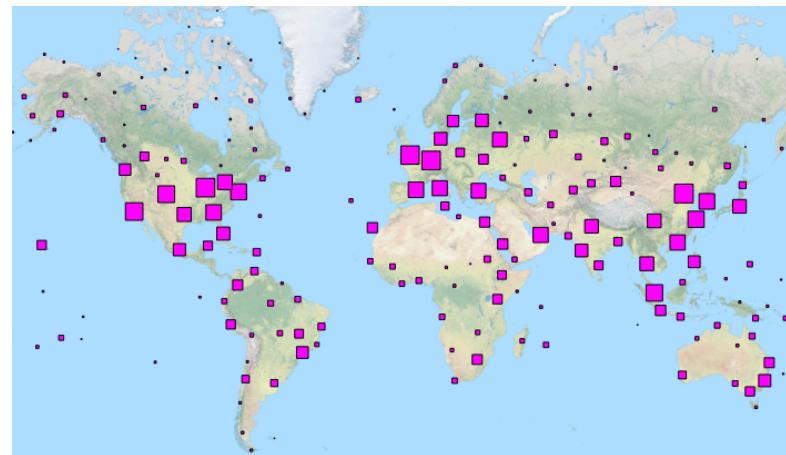
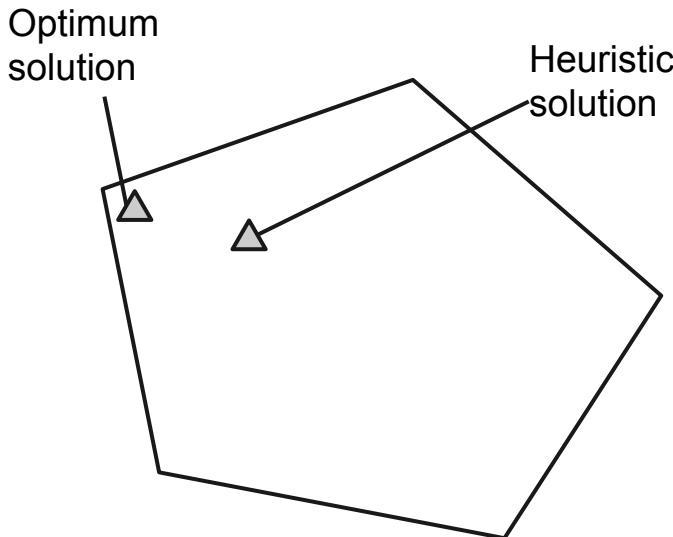
- Improving running time (in-memory, parallelism)
- Improve salience function (static weights too crude)
- Improve syntax (constraints are hard to define)
- Improve expressibility (e.g. relax zoom-consistency)



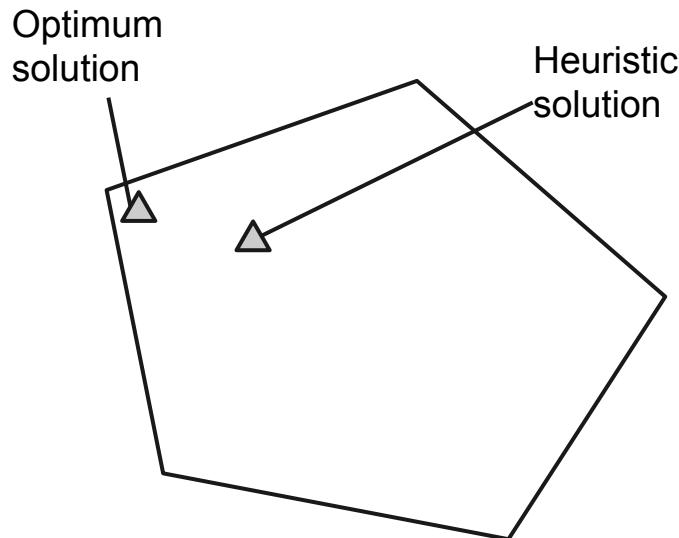
Thank you

BACKUP SLIDES

Algorithm matters!



Algorithm matters!



Which airport to keep?

